

Information loss of fixed linear health indicators as a control problem: simplex geometry and reachability limits

Huy Hoang Le¹ and Kim-Anh Nguyen²

¹R&D Department, PowerMore Ltd., Da Nang 550000, Viet Nam

²Faculty of Electrical Engineering, University of Science and Technology, The University of Danang, Da Nang 550000, Viet Nam

Correspondence: Kim-Anh Nguyen, Faculty of Electrical Engineering, University of Science and Technology, The University of Danang, Da Nang 550000, Viet Nam. Email: nkanh@dut.udn.vn

Email of the first author: lhhoang@powermore.vn

Running title: Health indicators as a control problem

Abstract

Health indicators combine condition-monitoring channels with fixed weights, yet the combination most informative about remaining life rotates as degradation mechanisms progress at different rates, so the indicator discards information. We treat that loss as a control problem: the operating point enters the nonlinear degradation dynamics and steers the informative direction. For a linear indicator, positive degradation rates and Gaussian remaining life over a time-shift population, the loss depends on the indicator only through one misalignment angle, a distance on a statistical simplex where degradation is an input-modulated replicator flow. Zero loss is attainable exactly when the informative directions reachable at each instant share a common point, which generically needs one fewer independent input than there are mechanisms. The best achievable loss is the value of a fractional optimal control problem whose terminal time, the failure time, is controlled; and certifying it reduces to one continuously differentiable verification function. No criterion built from the Fisher information about remaining life detects the loss. On a representative synthetic plant, steering lowers the information-loss objective 68-fold against the best constant operating point. Across 299 random plants of the separable model, the reachable-set width, known before design, provides a moderate screening signal for when steering pays (area under the receiver-operating curve 0.73). Public data show the rotation in measured battery cells and, against a negative control, support its dependence on the operating point in simulated turbofan records; the highest-fidelity turbofan data reject that separable form.

Keywords: fractional optimal control; reachable sets

information geometry; health indicator; prognostics; remaining useful life

1 Introduction

Condition monitoring rarely feeds raw measurements to a life-prediction algorithm. A vibration channel, a temperature, an impedance and a capacity reading first go into a single scalar, the health indicator, and the life-prediction algorithm sees only that scalar [1, 2]. The compression is almost always fixed: the designer picks one set of weights at design time, and every unit carries those weights for its whole life, because firmware embeds them, certification freezes them, or re-tuning them in the field is impractical.

Fixing the compression has a cost. Real components degrade through several mechanisms at once, and those mechanisms do not proceed in step: an early-life process may dominate for the first third of the record and then saturate, while a late-life process accelerates. The combination of channels that carries the most information about remaining life therefore rotates as the component ages. A fixed indicator is aligned with that combination at best at one moment, and is misaligned for the rest of the record; the information it discards is not recoverable downstream, however good the prediction algorithm that follows [3].

The established response is to make the map better: choose a different criterion for it, let it switch between segments of life, or learn it from more run-to-failure data [2, 4]. All of these take the degradation as given and work on the compression. This paper acts on the degradation itself, through the operating point.

The rotation is not a property of the instrumentation. The degradation physics produces it, and that physics responds to stress: temperature, load and duty cycle change the relative rates of the competing mechanisms, not only the overall speed. Running a cell hotter does not only age it faster; it also shifts the balance between the mechanisms, and therefore moves the direction in which information about remaining life arrives.

The operating point is, in most applications, some-

thing the operator chooses. So the misalignment between a fixed indicator and the informative direction is not fixed by nature. It is a function of the input, and the operator can steer it.

That reframing turns a data-analysis question into a control question, and control questions come with sharper answers than design heuristics do. One can ask what misalignments are reachable at a given operating envelope, whether zero loss is attainable at all, how many independent inputs that would take, and what the best achievable loss is when it is not. These are the questions this paper answers.

They sit against five bodies of work. Four address parts of the problem and the fifth supplies the geometry used in Section 3; to our knowledge none of them addresses the position taken here.

The prognostics literature treats the compression as a design problem to be solved from data: searching genetically over candidate features against monotonicity and trendability criteria [2], fusing the sensors linearly into a composite health index whose weights are fitted once to a degradation model [5] and analysing that fusion statistically [6], learning it with a recurrent network [7], or surveying the resulting pipelines [1, 8]. The composite indices of [5, 6] are fixed linear maps of exactly the kind studied here. Two limits on that programme are known, information-theoretic bounds on what any fixed compression can retain [3] and the number of run-to-failure units needed to learn one reliably [4]. The first is our own, and the boundary is worth drawing explicitly: it takes one degradation trajectory as given and asks how much information a fixed compression must lose along it. The question here is the complementary one, whether that trajectory can itself be steered so that the same fixed map loses less. Both prior limits hold the trajectory fixed; neither asks whether it can be changed.

The predictors downstream are overwhelmingly stochastic — Wiener, gamma and inverse-Gaussian degradation processes with first-passage failure [9, 10]; reviewed for remaining-life estimation in [11, 12] and embedded in system-level design methodology in [13]. That literature models how damage evolves, not how the measurement should be compressed before the model is applied, which is the question upstream of it.

A separate literature does treat the operating point as a decision variable, but to extend remaining life rather than to preserve estimability [14, 15]; the two aims can conflict, and Section 6 exhibits an instance in which they disagree. Choosing inputs to make a state easier to estimate is itself classical, going back to dual control [16], and its modern forms shape a scalarisation of the Fisher information matrix [17–19]. Our Proposition 4.6 shows the quantity at issue here to be invisible to every such scalarisation.

The fifth body is information geometry. The reformulation of Section 3 rests on the Fisher–Rao geometry of the probability simplex and on the Bhattacharyya coefficient [20, 21], and the flow it produces is the replicator

equation of evolutionary dynamics [22, 23]. These are standard objects; what is new here is their identification with the informative direction of a degrading component, and the fact that the fitness driving the flow is actuated.

The contributions are as follows.

1. The information a fixed indicator discards is an exact function of one misalignment angle (Lemma 2.2). Placing the channels’ shares of the Fisher information on a statistical simplex turns the degradation into a replicator flow whose rates the operating point modulates, and the optimal indicator into a barycentre of the degradation path (Section 3).
2. Zero loss is attainable exactly when the reachable direction sets share a point, which generically needs at least as many independent operating variables as mechanisms less one (Section 4). Supplying the missing input collapses the peak reachability floor from 4.52° to zero at every state, so the rule prices instrumentation rather than only bounding it. Taking logarithms makes every reachable set a translate of one fixed set, which turns the question into a containment and shows the counting rule to be sufficient but not necessary; the exceptions exist for any number of mechanisms and are too fragile to build on (Section 4.1). The count separates two objectives that are easily conflated: holding the misalignment at a constant level is a single scalar condition and generically needs one input, whereas freezing the informative direction outright is $d - 1$ conditions. That gap is the difference between useful steering and lossless steering, and an installation with one controllable variable can have the first without the second.
3. The loss is invariant to a common gain on all channels and is therefore invisible to every scalarisation of the Fisher information about remaining life, so no criterion built from that information alone can rank indicators by the loss they incur.
4. The best achievable loss is the value of an optimal control problem whose objective is a ratio, and whose terminal time is itself controlled (Section 5); we compute its value to two digits and do not certify it. The actuation authority, measured by the width of the reachable set and computable before any indicator is designed, provides a moderate screening signal for the regime in which steering can have substantial value; on 299 randomly drawn plants it ranks them as well as its ratio to the drift of the informative direction and fits them better, and the drift matters only in combination with it (Section 6.7).
5. The premises are separately testable on public run-to-failure data, and Section 7 tests three in turn:

Table 1: Where this paper sits. The columns are the axes that separate these literatures, not a ranking: each body of work answers a question the others do not ask, and the last column is a restriction this paper imposes on itself rather than an advantage over the others. “Input alters degradation” means the operating point enters the degradation dynamics; “direction” means the informative direction of Section 2 is itself an objective.

	input alters degradation	lifetime is an objective	direction is an objective	indicator fixed, scalar
Indicator design [2, 5–7]	—	—	—	✓
Fixed-indicator limits [3, 4]	—	—	—	✓
Stochastic life models [9, 11]	—	—	—	—
Health-aware control [14, 15]	✓	✓	—	—
Observability-aware [16, 17, 19]	✓	—	partly	—
This paper	✓	secondary	✓	✓

whether the informative direction rotates, whether the operating point moves it, and whether the separable stress model survives. The first holds on measured fleets, the second so far only in simulation, and the third is rejected. That section therefore reports which hypotheses withstand contact with data rather than demonstrating that the method works.

The results do not all rest on the same assumptions, and it is worth separating them before any of them is stated. Lemma 2.2, Proposition 2.3 and Theorem 4.1 need only the observation model of Section 2: a fixed linear indicator on whitened channels, and a scalar latent. Lemma 4.3 and Proposition 4.4 need more, namely that the input enters each mechanism through a single separable factor, and Section 7 reports that the highest-fidelity turbofan data reject exactly that. The empirical readings of Section 7 need more again, since estimating one informative direction from a fleet presumes the units differ by a time shift rather than by their degradation coefficients. Fig. 6 is the map, and the reader who wants to know which branch a given result sits on should consult it rather than infer it.

Section 2 states the model. Sections 3 to 5 develop the theory, Section 6 exercises the three of its claims that are stated sharply enough to be refuted, and Section 7 reports what public fleet data can and cannot establish about the premises the theory rests on.

2 Problem formulation

This section states the degradation model, says precisely what a prognostic algorithm is being asked to estimate, and where the uncertainty in that quantity comes from, derives the direction along which information about it arrives, and defines the fixed indicator whose information loss the rest of the paper studies. Readers familiar with [3] will recognise the model; the derivation is given in full here because the geometry of Section 3 depends on several of its details.

2.1 Degradation driven by a controlled operating point

A component degrades through d distinct mechanisms acting at once. In a lithium-ion cell these might be loss of cyclable lithium, growth of the solid-electrolyte interphase and mechanical fatigue of the electrode; in a rolling bearing, spalling, wear and lubricant degradation. We write $x_i(t) \geq 0$ for the damage accumulated by mechanism i up to time t , and collect these into the state $x = (x_1, \dots, x_d)$.

One externally chosen operating point drives all d mechanisms. It collects the temperature, the mechanical or electrical load, and the duty cycle at which the operator runs the component. We write that operating point as

$$u(t) \in \mathcal{U} \subset \mathbb{R}^m, \quad (1)$$

with $\mathcal{U}$ compact, representing the operating envelope the application permits. Throughout this paper u is the control input: it is the decision variable, and the question we ask is what choosing it well can buy. The number of independent inputs, m , turns out to decide what is achievable, and it is usually much smaller than the number of mechanisms d .

The damage states evolve according to

$$\dot{x}_i = \varphi_i(u) h_i(x_i), \quad i = 1, \dots, d, \quad (2)$$

in which the two factors play different roles:

- $\varphi_i(u) > 0$ is the stress-acceleration factor of mechanism i : how much faster that mechanism runs at operating point u than at a reference point. It carries all the dependence on the input.
- $h_i(x_i) > 0$ is the intrinsic shape of mechanism i : how its rate depends on the damage it has already done. It carries all the dependence on the state.

This separation is the standard accelerated-life form and covers the stress models in routine use [24]. In the Arrhenius model the input is absolute temperature and $\varphi_i(u) = A_i \exp(-E_i/(ku))$, with $A_i > 0$ a pre-exponential constant, $E_i > 0$ the activation energy of mechanism i and k Boltzmann’s constant. In the inverse-power-law model the input is a mechanical or electrical stress and $\varphi_i(u) = \kappa_i u^{n_i}$, with $\kappa_i > 0$ a

scale factor and $n_i > 0$ the stress exponent. Choosing h_i recovers the exponential, power-law and stretched-exponential degradation classes [3]; the results below use only $\varphi_i > 0$ and $h_i > 0$.

The within-unit dynamics are deterministic, where most of the degradation literature places the randomness inside a single unit's path [9, 10, 12]. We put it between units instead (Section 2.2). The object of study is the direction along which information arrives, which is fixed by the relative rates of the competing mechanisms: within-unit diffusion would jitter that direction without altering its systematic motion, whereas dispersion in the rate parameters across a population is what makes remaining life uncertain and therefore worth estimating.

We measure damage on a scale normalised so that mechanism i exhausts the component's tolerance when x_i reaches unity. A component is withdrawn at a condemnation limit at or below that point, so we write the end of life as the first passage of any one mechanism to a threshold $x_f \in (0, 1]$:

$$T_f = \inf\{t \geq 0 : \max_i x_i(t) \geq x_f\}. \quad (3)$$

No result below depends on the value of x_f , which enters only by fixing where a trajectory stops. The computed quantities of Section 6 do depend on it, and the value used there is stated with the rest of the test plan.

2.2 What is to be estimated, and why it is uncertain

The quantity a prognostic system must report is the remaining useful life

$$R(t) := T_f - t, \quad (4)$$

the time left before failure.

We must say where the uncertainty in R lies, because (2) is a deterministic differential equation: for one component whose initial damage and parameters were known exactly, T_f and hence R would be known exactly, and there would be nothing to estimate. The uncertainty is between components, not within one. Units built to the same specification differ in initial damage and in the coefficients of (2), which is the unit-to-unit variability that reliability engineering treats as the primary source of scatter in life data [24]. A prognostic algorithm observes one unit drawn from that population and must infer its remaining life.

We therefore treat R , at any given observation instant, as a random variable over the population of units, with prior variance $\text{Var}[R] > 0$. This variance is the uncertainty the algorithm exists to reduce, and it is the reason the estimation problem is not trivial. Everything below is written for a nominal trajectory of (2), with the fleet spread entering only through $\text{Var}[R]$.

2.3 Observations and the informative direction

Each damage state reaches the record through its own noisy channel,

$$y_i = x_i + \sigma_i \epsilon_i, \quad i = 1, \dots, d, \quad (5)$$

with $\sigma_i > 0$ the noise level of channel i and ϵ_i independent standard normal variables. The record available to the algorithm is the collection of these measurements over the observation window.

Equation (5) pairs each mechanism with a channel of its own. Condition monitoring in general returns $y = H(x, u) + \epsilon$, with sensors that mix mechanisms, and the pairing is an assumption rather than a description. It is load-bearing: the positivity of d that Remark 3.1 shows the results to need follows from $\varphi_i > 0$ and $h_i > 0$ in these coordinates and need not survive an arbitrary H . Where H is known, invertible and positivity-preserving the construction carries over to $H^{-1}y$; the general case does not reduce to it.

Two assumptions carry the sensitivity, and both are stated here rather than left to be inferred.

The prior on R We take R to be Gaussian over the population with variance $\text{Var}[R]$. Linearising the record about the nominal trajectory, following [3], then makes the record and R jointly Gaussian, and the mutual informations below are exact. Linearisation alone would not do it: with a non-Gaussian prior of the same variance the Gaussian value is an upper bound on what a channel carries, and Lemma 2.2 becomes an inequality rather than an identity. The information geometry that follows is exact under this Gaussian prior and the time-shift population described next; it is not a model of population variability in general.

How units differ We take the population spread to be a reparametrisation of a single degradation path: units differ in where they sit on the trajectory of (2), not in its shape. This is what makes R a coordinate along that path, and it is what the next display uses. Dispersion in the coefficients of (2) is a different matter, because it moves $\partial x / \partial \theta$ rather than sliding along $\partial x / \partial R$, and Remark 2.1 records what it costs.

The sensitivity of channel i to remaining life then follows from (4): because R falls at unit rate as t advances,

$$\frac{\partial x_i}{\partial R} = -\frac{\partial x_i}{\partial t} = -\varphi_i(u) h_i(x_i). \quad (6)$$

Remark 2.1 (what the time-shift assumption costs). *Equation (6) is exact when the population spread slides along one path, and only then. On the plant of Section 6.1, with units instead drawn by dispersing the degradation coefficients, the direction that the population regression of x on R actually returns stands about 45° from $-\varphi(u)h(x)$, across two families of dispersed coefficient and every spread tried; a control whose units*

are started on the reference trajectory, where (6) holds by construction, returns the same regression to within 0.2° . The misalignments this paper is about are a few degrees. Extending the construction to populations dispersed in θ , where the observation carries scatter that no value of R explains, is open.

Dividing each channel by its own noise level puts them on a common scale, and taking magnitudes gives the whitened signal direction

$$\mathbf{d}(t) = [\sigma_i^{-1} \varphi_i(u) h_i(x_i)]_{i=1}^d, \quad \hat{\mathbf{d}} = \mathbf{d} / \|\mathbf{d}\|. \quad (7)$$

In words: $\mathbf{d}$ says how strongly, and in what combination of channels, the measurement responds to a change in remaining life. Its direction $\hat{\mathbf{d}}$ is the *informative direction*, the combination a health indicator has to be aligned with and the object named informally in Section 1. Its length $\|\mathbf{d}\|$ is the strength of that signal. The two objects carry the paper: $\|\mathbf{d}\|$ is how much information arrives and $\hat{\mathbf{d}}$ is where it arrives from, and the loss of a fixed indicator depends only on the second. Experiment design in the observability tradition scores the first, which is why Proposition 4.6 finds it blind to the loss.

We use two properties of (7) repeatedly below.

Positivity Every entry of $\mathbf{d}$ is strictly positive, because $\varphi_i > 0$ and $h_i > 0$ in (2). This is a consequence of the degradation model rather than an assumption imposed for convenience, and Remark 3.1 shows that the results below would fail without it.

Actuation $\mathbf{d}$ is an explicit algebraic function of the state and the input. The direction along which information about R arrives is therefore not a fixed property of the instrumentation: it moves when the operating point moves. That is the observation this paper is built on, and Fig. 1 places it: the operating point reaches $\hat{\mathbf{d}}$ twice, once through the state and once directly.

Because R is scalar, a single signal-to-noise ratio governs the information the record carries about it, namely the squared sensitivity times the prior uncertainty,

$$\gamma = \|\mathbf{d}\|^2 \text{Var}[R]. \quad (8)$$

Rates in (7) are taken with respect to clock time, whereas [3] differentiates with respect to age normalised by T_f . The two conventions differ by a factor common to all d channels, so they give the same $\hat{\mathbf{d}}$, and hence the same geometry in Sections 3–4, while rescaling γ . Section 5, where lifetime enters the objective, restores the normalised convention.

2.4 The fixed health indicator and its information loss

A health indicator compresses the d measured channels into one scalar before any life prediction is made.

We take it to be a fixed linear map. Since (7) already measures each channel in units of its own noise, we write the indicator in those same coordinates: with $\Sigma := \text{diag}(\sigma_1^2, \dots, \sigma_d^2)$ and $\tilde{y} := \Sigma^{-1/2} y$ the whitened record,

$$z = \mathbf{v}^\top \tilde{y}, \quad \|\mathbf{v}\| = 1, \quad (9)$$

with the weight vector $\mathbf{v}$ chosen once, at design time, and held for the whole life of the component. Fixing it is not a modelling convenience: the map is typically embedded in firmware or frozen by certification, so it cannot be re-tuned as the component ages. An indicator specified on the raw measurements, $z = \mathbf{w}^\top y$, is the same map under $\mathbf{v} \propto \Sigma^{1/2} \mathbf{w}$, and the normalisation $\|\mathbf{v}\| = 1$ only fixes the scale of z ; it restricts nothing. The distinction matters because the angle below is measured after whitening: unless the σ_i are equal, a unit vector in raw coordinates is not a unit vector in whitened ones, and the two readings of ψ disagree.

We compute the cost of fixing $\mathbf{v}$ as follows. Write $\psi(t)$ for the angle between the informative direction $\hat{\mathbf{d}}(t)$ and the indicator weights $\mathbf{v}$. After whitening, the signal in the record points along $\hat{\mathbf{d}}$ and has magnitude $\|\mathbf{d}\| \sqrt{\text{Var}[R]}$, while the noise is isotropic of unit scale. Projecting onto a unit vector $\mathbf{v}$ therefore keeps the fraction $\cos \psi$ of the signal and all of the noise, so the indicator's signal-to-noise ratio is

$$\gamma \cos^2 \psi \quad (10)$$

against the γ of (8) available to an oracle that is free to re-aim its projection at $\hat{\mathbf{d}}(t)$ at every instant. Since a scalar Gaussian channel of signal-to-noise ratio s carries $\frac{1}{2} \log(1+s)$ nats about R [25], the information the fixed indicator forfeits, at each instant, is


$$\ell = \frac{1}{2} [\log(1+\gamma) - \log(1+\gamma \cos^2 \psi)] \geq 0, \quad (11)$$

The design objective throughout is its average over the life of the component,

$$\bar{\ell} := \frac{1}{T_f} \int_0^{T_f} \ell(t) dt, \quad (12)$$

a loss per unit time. We optimise this rate rather than the information the whole observation record forfeits, because (12) is unaffected by how densely the record is sampled and is therefore a property of the indicator rather than of the measurement schedule. The two are different quantities, and Proposition 2.3 gives the second in closed form.

A third quantity is the one an operator actually pays for: the error of whatever remaining-life estimator runs downstream. It is neither of the first two and is not a function of them. The loss is a logarithmic information quantity and the error is a time, and an estimator holding a whole history recovers much of what any single instant discards, so a large spread in (12) does not imply a comparable spread in error. What the theory can offer downstream is therefore an ordering rather than a conversion, and Section 6.8 tests that ordering against an estimator built from none of this construction: it survives,



Alignment asks $\hat{\mathbf{d}}(x, u) = q$: $d - 1$ scalar equations on m unknowns, whence the counting rule of Theorem 4.1.

Figure 1: The construction. The operating point enters twice, through the plant and directly through $\varphi_i(u)$; the second path is what makes the misalignment ψ a control variable. One input drives all d mechanisms.

with a spread in error much smaller than the spread in loss. Equation (12) should be read accordingly, as an information-geometric objective whose ranking carries downstream, not as a surrogate for prognostic accuracy.

Two remarks on what (11) does and does not say. It is a loss relative to an oracle that no fixed map can realise, so it measures the price of the time-invariance constraint rather than of a bad design; a perfectly chosen $\mathbf{v}$ still pays it whenever $\hat{\mathbf{d}}$ moves. And it depends on the indicator only through the single angle ψ : two indicators that make the same angle with $\hat{\mathbf{d}}$ are equally good, whatever else distinguishes them.

The next lemma replaces (11) by a closed form that will be more convenient throughout, and records how it behaves at small misalignment.

Lemma 2.2. *Let $w := \gamma/(1 + \gamma) \in (0, 1)$. Then*

$$\ell = -\frac{1}{2} \log(1 - w \sin^2 \psi), \quad (13)$$

exactly and at every signal-to-noise ratio, and

$$\ell = \frac{w}{2} \psi^2 + \frac{w(3w-2)}{12} \psi^4 + O(\psi^6). \quad (14)$$

Proof. Substituting $\cos^2 \psi = 1 - \sin^2 \psi$ into (11) gives

$$\ell = -\frac{1}{2} \log \frac{1 + \gamma - \gamma \sin^2 \psi}{1 + \gamma} = -\frac{1}{2} \log(1 - w \sin^2 \psi).$$

For (14), use $-\frac{1}{2} \log(1 - s) = \frac{s}{2} + \frac{s^2}{4} + O(s^3)$ with $s = w \sin^2 \psi$ and $\sin^2 \psi = \psi^2 - \psi^4/3 + O(\psi^6)$; collecting the ψ^4 terms gives $-w/6 + w^2/4 = w(3w - 2)/12$.

The quadratic term of (14) is the bound obtained in [3, Thm. 1], which is therefore the leading term of an exact identity rather than an approximation of unknown quality. The quartic coefficient $w(3w - 2)/12$ changes

sign at $w = 2/3$, that is at $\gamma = 2$: below that signal-to-noise ratio the quadratic model overstates the loss, above it the model understates it. This matters in practice because prognostic problems sit on both sides of $\gamma = 2$: an early-life measurement with a weak signal on one side, a well-instrumented component near failure on the other. Fig. 2(a) plots the exact loss against its quadratic approximation on both sides of that ratio.

Proposition 2.3. *Let the record consist of measurements (5) at instants $t_1, \dots, t_N$, with the noise independent across channels and across instants, and write $\mathbf{d}_k := \mathbf{d}(t_k)$ and ψ_k for the angle between $\hat{\mathbf{d}}_k$ and $\mathbf{v}$. Put*

$$G := \text{Var}[R] \sum_{k=1}^N \|\mathbf{d}_k\|^2, \quad W := \frac{G}{1 + G},$$

$$S := \frac{\sum_k \|\mathbf{d}_k\|^2 \sin^2 \psi_k}{\sum_k \|\mathbf{d}_k\|^2}.$$

Then the information about R that the fixed indicator forfeits over the whole record is

$$\ell_{\text{rec}} = -\frac{1}{2} \log(1 - WS). \quad (15)$$

Proof. See Supplementary Section S2. $\square$

Equation (15) is (13) again, with the instantaneous signal-to-noise ratio γ replaced by its sum G over the record, and $\sin^2 \psi$ by the mean S weighted by the signal energy $\|\mathbf{d}_k\|^2$. The reduction to one angle therefore survives the passage from an instant to a record. What does not survive is the averaging: $\bar{\ell}$ averages the loss, whereas (15) averages $\sin^2 \psi$ inside it, and the two differ. They also select different indicators. Minimising (15) over $\|\mathbf{v}\| = 1$ maximises $\mathbf{v}^T (\sum_k \mathbf{d}_k \mathbf{d}_k^T) \mathbf{v}$, so the




Figure 2: (a) Lemma 2.2: the exact loss (solid) against its quadratic model (dotted), whose quartic coefficient changes sign at $\gamma = 2$. (b) Theorem 3.6: of the replicator velocity, only the component along the geodesic towards q moves β . (c) The rotation, measured on the Severson fleet at two differentiation windows.

best fixed indicator for the record is the leading eigenvector of the integrated Fisher matrix, while the minimiser of $\bar{\ell}$ is the barycentre of Proposition 3.3. Every result below is stated for $\bar{\ell}$.

Corollary 2.4 (what the loss costs an estimator). *In the model of Section 2, the minimum-variance estimate of R formed from the fixed indicator carries mean-squared error*

$$\frac{\text{Var}[R | z]}{\text{Var}[R | \tilde{y}]} = \frac{1 + \gamma}{1 + \gamma \cos^2 \psi} = e^{2\ell}$$

times that of the estimate formed from the whole whitened record, at each instant. The same identity holds over a record with ℓ replaced by ℓ_{rec} of Proposition 2.3.

Proof. Write $\tilde{y} = \mathbf{d}R + \epsilon$ with $\epsilon \sim N(0, I_d)$ and R Gaussian. The indicator returns $z = (\mathbf{v}^\top \mathbf{d})R + \mathbf{v}^\top \epsilon$ whose noise is standard normal because $\|\mathbf{v}\| = 1$, so its signal-to-noise ratio is $\gamma \cos^2 \psi$ and $\text{Var}[R | z] = \text{Var}[R]/(1 + \gamma \cos^2 \psi)$. The record’s sufficient statistic is $\mathbf{d}^\top \tilde{y}$, with ratio γ , giving $\text{Var}[R | \tilde{y}] = \text{Var}[R]/(1 + \gamma)$. Their quotient is $e^{2\ell}$ by (11). The record statement is the same computation with G and S in place of γ and $\sin^2 \psi$.

The loss is therefore not a score standing in for estimation error: it is the logarithm of the factor by which that error inflates. What the corollary does not settle is whether the ordering survives an estimator other than the one the model prescribes, and Section 6.8 checks that separately.

3 The information simplex and its replicator flow

Lemma 2.2 reduced the loss to one angle. This section shows that the angle lives on a familiar object, the probability simplex, and that the degradation traces a flow along it that the operating point drives. The payoff is threefold: the loss becomes a recognised statistical distance, the optimal indicator becomes a barycentre, and the effect of the input becomes a single interpretable

quantity. Two of the results here are load-bearing for what follows — the replicator form of Theorem 3.4 and the misalignment rate of Theorem 3.6, which Section 4 builds the reachability argument on. The Bhattacharyya identity and the barycentre are interpretation: they say what the loss and the optimal indicator already are in recognised terms, and nothing later depends on them.

3.1 Information shares

Because every entry of $\mathbf{d}$ is positive (Section 2.3), the squared components of the unit vector $\hat{\mathbf{d}}$ form a probability vector:

$$p_i := \hat{\mathbf{d}}_i^2, \quad p_i > 0, \quad \sum_i p_i = 1. \quad (16)$$

The interpretation is direct. The total Fisher information the record carries about R is $\|\mathbf{d}\|^2 \text{Var}[R]$ by (8), and channel i contributes $\mathbf{d}_i^2 \text{Var}[R]$ of it. So p_i is exactly the share of the total information about remaining life that arrives through channel i . A component in which one mechanism dominates the degradation has p near a vertex of the simplex; one in which three mechanisms contribute equally sits near its centre.

The same map applied to the indicator, $q_i := \mathbf{v}_i^2$, is legitimate because an optimal indicator may be taken with non-negative weights.

Remark 3.1 (positivity is load-bearing). *If $\hat{\mathbf{d}} \geq 0$ then $\hat{\mathbf{d}}^\top |\mathbf{v}| \geq |\hat{\mathbf{d}}^\top \mathbf{v}|$ at every instant, so replacing $\mathbf{v}$ by $|\mathbf{v}|$ never increases (11): non-negative weights are without loss of generality. This argument uses the positivity of $\mathbf{d}$ inherited from $\varphi_i, h_i > 0$, and it fails without it: in random trials with sign-indefinite $\hat{\mathbf{d}}$, taking absolute values was worse than the signed vector about a third of the time. The results therefore require the positive-orthant structure of the degradation model; it is not only a notational convenience.*

3.2 The loss as a Bhattacharyya distance

The identification below is used twice in what follows: it gives “the average of a curve on the simplex” an unambiguous meaning (Section 3.3), and it lets the high-signal limit of the loss be read off as a known divergence.

Proposition 3.2. $\cos \psi = \sum_i \sqrt{p_i q_i}$ is the Bhattacharyya coefficient of p and q . Consequently $\psi = \frac{1}{2} d_{\text{FR}}(p, q)$, where d_{FR} is the Fisher–Rao distance on the simplex, and, writing $D_{\text{B}}(p, q) = -\log \sum_i \sqrt{p_i q_i}$ for the Bhattacharyya divergence,

$$\ell = -\frac{1}{2} \log[(1 - w) + w e^{-2D_{\text{B}}(p, q)}]. \quad (17)$$

As $\gamma \rightarrow \infty$, $\ell \rightarrow D_{\text{B}}(p, q)$.

Proof. By construction,

$$\cos \psi = \hat{\mathbf{d}}^\top \mathbf{v} = \sum_i \sqrt{p_i} \sqrt{q_i},$$

which is precisely the Bhattacharyya coefficient of p and q . The map $p \mapsto \sqrt{p}$ is the classical isometry between the simplex under the Fisher–Rao metric and the positive orthant of the unit sphere, under which the geodesic distance is twice the spherical angle [20, 21]; hence $\psi = \frac{1}{2}d_{\text{FR}}$. Then $\cos^2 \psi = e^{-2D_{\text{B}}}$, and substituting into (13) gives (17). Letting $w \rightarrow 1$ gives the limit.

The square root is a bijection between the simplex and the non-negative part of the unit sphere, carrying p to $\hat{\mathbf{d}}$ and q to $\mathbf{v}$, and by the proposition it is an isometry. We use the two descriptions interchangeably from here on: a statement such as $q \in \mathcal{D}(x)$ in Section 4 says that $\mathbf{v}$ is a reachable direction, and a distance between points of the simplex is the angle between the corresponding unit vectors.

The resulting picture is the one the rest of the paper develops. The degrading component is a curve on the simplex, the health indicator is a single point on it, and the information lost depends only on the distance between them. Design of the indicator is the choice of one point; control of the operating point is the shaping of the curve.

3.3 The optimal indicator as a barycentre

If the indicator is one point and the plant is a curve, the best point should be some kind of average of the curve. It is, but of two different kinds in two different limits.

Proposition 3.3. *Let q^* minimise the time-averaged loss.*

(i) *As $\gamma \rightarrow \infty$, $\ell \rightarrow D_{\text{B}}$ pointwise, q^* tends to the Bhattacharyya barycentre $\arg \min_q \int D_{\text{B}}(p_\tau, q) d\tau$.*

(ii) *As the curve shortens, $\ell = \frac{w}{2}\psi^2 + O(\psi^4)$, so q^* tends to the w -weighted Fisher–Rao barycentre, and $\bar{\ell} = \frac{1}{3}w \text{Var}^{\text{FR}} + O(\psi^4)$ where $\text{Var}^{\text{FR}}[\min_q \mathbb{E}_\mu[d_{\text{FR}}^2(p_\tau, q)]$ is the μ -weighted Fréchet variance of the curve.*

Proof. See Supplementary Section S2.

The two limits are taken in different parameters, signal-to-noise ratio in (i) and arc length in (ii), and they do not coincide. On a curve of angular span 82.7° the two barycentres lie 0.72° apart, and q^* crosses from the Fisher–Rao side to the Bhattacharyya side between 0 and 10 dB, reaching 3×10^{-5} degrees of the latter by 50 dB. For design purposes the practical consequence is mild: either barycentre of a nominal degradation path is a good starting indicator.

3.4 The replicator flow and the role of the input

We now ask how fast the informative direction moves and what the operating point does to that motion.

Theorem 3.4. *Let $g_i = \log \mathbf{d}_i$ denote the logarithmic rate of channel i . Along (2),*

$$\dot{\hat{\mathbf{d}}}_i = \hat{\mathbf{d}}_i(\dot{g}_i - \mathbb{E}_p[\dot{g}]), \quad (18)$$

so the information shares obey the replicator equation of evolutionary dynamics, in which each share grows in proportion to the amount by which its own rate exceeds the average, the average being taken over the shares themselves:

$$\dot{p}_i = 2p_i(\dot{g}_i - \mathbb{E}_p[\dot{g}]), \quad (19)$$

and

$$\|\dot{\hat{\mathbf{d}}}\| = \text{Std}_p[\dot{g}], \quad \dot{g} = [\varphi_i(u)h'_i(x_i)]_i + \nabla_u \log \varphi \dot{u}. \quad (20)$$

Proof. From $\mathbf{d}_i = e^{g_i}$ we have $\dot{\mathbf{d}}_i = \mathbf{d}_i \dot{g}_i$, so $\dot{\hat{\mathbf{d}}} = \|\mathbf{d}\|^{-1}(\dot{\mathbf{d}} - \hat{\mathbf{d}}(\hat{\mathbf{d}}^\top \dot{\mathbf{d}}))$ has components $\hat{\mathbf{d}}_i \dot{g}_i - \hat{\mathbf{d}}_i \sum_j \hat{\mathbf{d}}_j^2 \dot{g}_j$, which is (18); squaring and summing gives $\text{Var}_p[\dot{g}]$. Since $p_i = \hat{\mathbf{d}}_i^2$, $\dot{p}_i = 2\hat{\mathbf{d}}_i \dot{\hat{\mathbf{d}}}_i = 2p_i(\dot{g}_i - \mathbb{E}_p[\dot{g}])$. Differentiating $g_i = \log \varphi_i(u) + \log h_i(x_i) - \log \sigma_i$ along (2) gives $\dot{g}_i = \nabla_u \log \varphi_i \cdot \dot{u} + (h'_i/h_i)\varphi_i h_i$, and $(h'_i/h_i)\varphi_i h_i = \varphi_i h'_i$. $\square$

Equation (20) has a compact reading. The speed at which the informative direction rotates equals the information-weighted spread of the mechanisms' logarithmic rates. If all mechanisms accelerate at the same proportional rate the direction does not move at all, however fast the component is degrading; it is disagreement among the mechanisms that rotates it. The identity is the Fisher–Rao/Shahshahani geometry of the simplex combined with Fisher's fundamental theorem of natural selection [20, 22, 23]. What is new here is its identification with the informative direction of a degrading component, and, decisively, that the quantity playing the role of fitness, the logarithmic rate $\dot{g}_i$ of channel i , is actuated: the second summand in (20) is the operator's lever.

Remark 3.5 (contrast with Oja's flow). *Oja's flow for online principal-component extraction has the same algebraic form; for a diagonal covariance it reduces to a replicator equation with fitness twice the eigenvalues [26]. The roles are opposite. There the flow is an algorithm, run to fixation in order to find a direction, with constant fitness. Here it is the plant, its fitness is time-varying and actuated, and the aim is to prevent fixation from carrying the direction away from a fixed indicator.*

Only one scalar matters for the loss. Fix the indicator q and write $\beta := \angle(\hat{\mathbf{d}}, \mathbf{v}) = \frac{1}{2}d_{\text{FR}}(p, q)$ for the misalignment. This is the angle ψ of Lemma 2.2 under a new name, and the change of name marks a change of role: ψ was a fixed angle between two directions, whereas $\beta = \beta(x_t, u_t)$ is followed along the trajectory and is the quantity the input acts on. Every loss expression below may therefore be read as a function of β .

Theorem 3.6. *Let the operating-point trajectory $u(\cdot)$ be absolutely continuous, and let $\hat{e}$ be the unit tangent at $\hat{\mathbf{d}}$ along the geodesic towards $\mathbf{v}$. Then $\dot{\beta} = -\langle \hat{\mathbf{d}}, \hat{e} \rangle$, which is affine in the input rate: $\dot{\beta} = \alpha_0(x, u) + \alpha(x, u)^\top \dot{u}$. In particular $|\dot{\beta}| \leq \|\hat{\mathbf{d}}\|$, with equality only for motion along that geodesic.*

Proof. $\cos \beta = \hat{\mathbf{d}}^\top \mathbf{v}$ gives $-\sin \beta \dot{\beta} = \hat{\mathbf{d}}^\top \dot{\mathbf{v}}$. Also $\hat{e} = (\mathbf{v} - \cos \beta \hat{\mathbf{d}}) / \sin \beta$ and $\hat{\mathbf{d}}^\top \hat{\mathbf{d}} = 0$, so $\langle \hat{\mathbf{d}}, \hat{e} \rangle = \hat{\mathbf{d}}^\top \mathbf{v} / \sin \beta = -\dot{\beta}$. Affinity follows because $\dot{g}$ is affine in $\dot{u}$ by (20) and $\hat{\mathbf{d}}$ is linear in $\dot{g}$.

The hypothesis is needed only to write $\dot{u}$; nothing in Section 4 depends on it, since the reachable sets (21) are defined pointwise in u and the limits below hold for every measurable admissible input. Theorem 3.6 is a local kinematic description of smooth steering, not a restriction on the controls the rest of the paper admits.

The component of the replicator velocity orthogonal to $\hat{e}$ is therefore free: the informative direction may rotate in a $(d-2)$ -dimensional subspace at no cost in information. Fig. 2(b) shows the decomposition at one state. This is the observation that separates the two objectives of the next section.

4 Reachability limits on the misalignment

Controlling the rate $\dot{\beta}$ is not the same as reaching a value of β . The distinction is the subject of this section, and it is what decides whether a fixed indicator can be made lossless at all.

For a state x , the set of informative directions the operator can produce by choosing the operating point is

$$\mathcal{D}(x) := \{\hat{\mathbf{d}}(x, u) : u \in \mathcal{U}\} \subset S^{d-1}, \quad (21)$$

the reachable direction set, and the closest the indicator can be brought at that state is the reachable misalignment floor

$$\beta_{\min}(x, q) := \text{dist}(q, \mathcal{D}(x)). \quad (22)$$

For a single input, $\mathcal{D}(x)$ is an arc traced out as u sweeps $\mathcal{U}$; its length measures the actuation authority available at x , and it shrinks to a point when the mechanisms respond identically to the input.

Theorem 4.1. *Fix an admissible input and let $x(\cdot)$ be the trajectory it generates. Zero loss along that trajectory is attainable with the fixed indicator q if and only if $q \in \bigcap_t \mathcal{D}(x_t)$. If $m < d-1$ and $u \mapsto \hat{\mathbf{d}}(x, u)$ is continuously differentiable, the indicators that attain zero loss form a null subset of the unit sphere, so almost every q , and in fact every q that does not attain it, has $\beta_{\min} > 0$ on a set of times of positive measure. Whether any q at all attains zero loss is a property of the plant, and Remark 4.5 exhibits plants for which one does. In every case $\bar{\ell} \geq \int \ell(\beta_{\min}(x_t, q)) dt / T_f$ for that trajectory.*

Proof. By Lemma 2.2, $\ell(\beta) = 0$ iff $\sin \beta = 0$, so $\bar{\ell} = 0$ iff $\beta_t = 0$ for almost every t , that is iff $\hat{\mathbf{d}}(x_t, u_t) = q$, that is iff $q \in \mathcal{D}(x_t)$ for every t . Solving $\hat{\mathbf{d}}(x, u) = q$ imposes $d-1$ independent scalar equations, both sides being unit vectors, on m unknowns. If $m < d-1$ then $\mathcal{D}(x)$, the image of an m -dimensional set under a continuously differentiable map, is a null subset of the $(d-1)$ -dimensional sphere, which almost every q misses; if $m \geq d-1$ the implicit function theorem supplies a solution wherever the Jacobian has full rank. The inequality $\beta_t \geq \beta_{\min}(x_t, q)$ holds for every admissible input by definition of $\beta_{\min}$ as a minimum over $\mathcal{U}$, and is strict exactly where $q \notin \mathcal{D}(x_t)$. Since $\mathcal{U}$ is compact and $\hat{\mathbf{d}}$ continuous, each $\mathcal{D}(x_t)$ is compact and varies continuously with t , so the set of times at which $q \notin \mathcal{D}(x_t)$ is open. It is non-empty for every q that does not attain zero loss, in particular for every q outside the null set $\mathcal{D}(x_0)$, and so has positive measure. Since ℓ is increasing on $[0, \pi/2]$, the bound follows on dividing by T_f . $\square$

This is a state-conditioned result. The bound is genuine, in the sense that no input violates a reachability constraint at the state it occupies, but it is conditional on the trajectory, which the input also steers. Optimising over the trajectory as well is deferred to Section 5, and the two levels are the structure of the argument: a state-conditioned limit here, a trajectory-level limit there.

The practical reading of Theorem 4.1 is a counting rule: a component degrading through d mechanisms generically needs $d-1$ independently adjustable operating variables before a fixed indicator can be kept perfectly aligned. Here *generically* is a count of dimensions over the plant's data, not a statement of measure: the theorem proves the measure statement over indicators, Remark 4.5 constructs plants that escape the count, and we do not prove those plants rare. Section 4.1 says what the qualifier covers. For a generic plant under this model, three mechanisms and a single temperature set-point will not do it. The three results of this section answer different questions, and it is worth separating them. Theorem 4.1 characterises zero loss exactly, as membership of every reachable set, and then counts dimensions to say when a generic plant can meet it; the count is necessary for generic data and is not by itself a construction. Proposition 4.4 gives an exact criterion, necessary and sufficient, on the separable subclass alone. Remark 4.5 exhibits plants outside that subclass which attain zero loss with fewer inputs than the count asks, so the count is not necessary in general. But a weaker and much cheaper objective is available.

Corollary 4.2 (what actuation is actually required). *Two objectives that are easily conflated carry different requirements.*

- (i) *Driving $\bar{\ell}$ to zero is equivalent to freezing $\hat{\mathbf{d}}$ at q , and needs $d-1$ conditions on $\dot{u}$, hence generically $m \geq d-1$.*

(ii) Holding β at a constant level, not necessarily zero, is the single scalar equation $\alpha_0 + \alpha^\top \dot{u} = 0$, solvable whenever $\alpha \neq 0$ and the resulting $u(\cdot)$ stays in $\mathcal{U}$; generically $m \geq 1$ suffices.

Proof. (i) By Theorem 4.1, $\bar{\ell} = 0$ iff $\hat{\mathbf{d}}(x_t, u_t) = q$ for every t , so $\hat{\mathbf{d}}$ is constant; by Theorem 3.4 that forces $\dot{g}_i = \mathbb{E}_p[\dot{g}]$ for every i , which is $d - 1$ independent conditions on $\dot{u} \in \mathbb{R}^m$. (ii) is one equation in $\dot{u}$ by Theorem 3.6. The gap between them is the content of that theorem: $\hat{\beta}$ sees only the component of $\hat{\mathbf{d}}$ along $\hat{e}$, so the remaining $d - 2$ directions of rotation are unconstrained by (ii) and constrained by (i).

Minimising $\bar{\ell}$ is neither of these: it trades the level of β against the trajectory that carries it, and is the problem of Section 5.

4.1 The reachable sets are translates of one set

Theorem 4.1 counts equations against unknowns and hedges with “generic”. The hedge is not idle, and the following observation says exactly what it covers.

Lemma 4.3 (separability). *Write $\xi(x, u)$ for the class of $\log \mathbf{d}(x, u)$ modulo the all-ones vector, which is what $\hat{\mathbf{d}}$ retains. Then*

$$\xi(x, u) = a(x) + b(u),$$

$$a_i(x) = \log \frac{h_i(x_i)}{\sigma_i}, \quad b_i(u) = \log \varphi_i(u), \quad (23)$$

both taken modulo the all-ones vector. Consequently every reachable direction set is the same set translated: $\mathcal{D}(x) = a(x) + \mathcal{B}$ with $\mathcal{B} := b(\mathcal{U})$ fixed.

Proof. Immediate from $\mathbf{d}_i = \sigma_i^{-1} \varphi_i(u) h_i(x_i)$ in (7): the logarithm turns the product into a sum, and the normalisation of $\hat{\mathbf{d}}$ removes any component along the all-ones vector.

The decomposition separates the two parties. The set $\mathcal{B}$ is fixed by the stress model φ and the envelope $\mathcal{U}$ alone (Fig. 3(a)). An operator’s authority can therefore be computed offline from the accelerated-life model, with no degradation record at all. The curve $\mathcal{A} := \{a(x_t)\}$ is fixed by the damage shapes h and the noise levels σ , and every reachable set is that one set carried along it (Fig. 3(b)). What the geometry of Section 3 measures is how far the plant’s demand $\mathcal{A}$ sticks out of the operator’s authority $\mathcal{B}$.

That offline route is the most useful consequence drawn in this paper and also the most restrictive. It holds on the separable subclass and not beyond it, and the best high-fidelity turbofan data we have reject separability (Section 7). This subsection and Proposition 4.4 should therefore be read as describing a structured subclass that is convenient when it applies, not as a general account of real machinery. Fig. 6 marks the one branch of the development that depends on it.

Proposition 4.4 (exact criterion). *Zero loss is attainable along a trajectory if and only if $\mathcal{A}$ is contained in a translate of $-\mathcal{B}$. In particular $\mathcal{A}$ must lie in an affine subspace of dimension at most $\dim \mathcal{B} \leq m$, which for a curve in the $(d - 1)$ -dimensional quotient fails generically unless $m \geq d - 1$; this recovers Theorem 4.1. When it does not fail, $m < d - 1$ suffices.*

Proof. By Theorem 4.1, zero loss holds iff $q \in \mathcal{D}(x_t)$ for every t . Writing κ for the quotient class of $\log \mathbf{v}$ and using Lemma 4.3, that reads $\kappa - a(x_t) \in \mathcal{B}$ for every t , that is $\mathcal{A} \subseteq \kappa - \mathcal{B}$. The dimension statement follows because a translate of $\mathcal{B}$ spans an affine subspace of dimension $\dim \mathcal{B}$. $\square$

The containment is drawn in Fig. 3(c), and the mechanism the count obscures is now visible. Alignment at an instant fixes u_t outright, since $b(u_t) = \kappa - a(x_t)$ determines it; the input is spent, and any residual constraint on the state must then hold along the flow with no freedom left to enforce it. One input can carry the whole burden only when no residual constraint is left, which happens when $\mathcal{A}$ already lies parallel to $\mathcal{B}$.

Remark 4.5 (a family with $m = 1$ and any d). *Take exponential shapes $h_i(x) = e^{\nu_i(x_i - x_i^0)}$, which (2) admits, and Arrhenius factors. Under the ramp $u(t) = u_0 + rt$ the damage rates are constant, $x_i(t) = x_i^0 + \varphi_i^0 t$ so $\varphi_i^0 = \varphi_i(u_0)$, so $a(x_t)$ moves at constant velocity $\nu_i \varphi_i^0$. Choosing $\nu_i = r E_i / \varphi_i^0$ makes that velocity parallel to E , which is the direction of $\mathcal{B}$, and $b(u_t)$ retreats along $\mathcal{B}$ at the same rate. The two cancel and $\hat{\mathbf{d}}$ is frozen. Supplementary Section S4 integrates the construction for $d = 2$ to 8.*

This bears on the sharpness of Theorem 4.1. The family of Remark 4.5 is a knife edge. At one percent of error in its shape coefficients, steering is already worse than not steering (Supplementary Section S4). A designer should read Proposition 4.4 as an explanation of when the counting rule binds, and Theorem 4.1 as the rule to design against.

4.2 Invisibility to Fisher-information criteria

A natural objection is that all of this should follow from designing the input to maximise the Fisher information about R , as in the observability-aware literature discussed in Section 1. It does not, and the reason is structural.

Proposition 4.6 (Fisher blindness). *The loss of a fixed indicator is not a function of the Fisher information about R : configurations with equal Fisher information can have different losses, and configurations with equal losses can have different Fisher information. No criterion depending on the Fisher information alone can therefore rank fixed indicators by the loss they incur.*




Figure 3: How the counting rule becomes a containment. (a) The operator's authority $\mathcal{B} = b(\mathcal{U})$, fixed by the stress model and the envelope alone. (b) Every reachable set is that one set translated, $\mathcal{D}(x_t) = a(x_t) + \mathcal{B}$ (circle: start of life, square: failure). (c) Zero loss asks the whole curve to lie in one translate (Proposition 4.4).

Proof. For a scalar latent the Fisher information is the number $\|\mathbf{d}\|^2 \text{Var}[R]$, so trace, determinant and least eigenvalue coincide and every scalarisation is a strictly monotone function of it. Under the rescaling $\mathbf{d} \mapsto c\mathbf{d}$ with $c > 0$, that number changes while $\hat{\mathbf{d}}$, and hence ℓ and ℓ , do not. A common gain therefore equalises the scores of two configurations with different losses, and separates the scores of two configurations with the same loss. $\square$

Stated in engineering terms, a Fisher criterion measures the strength of the signal, whereas the loss of a fixed indicator depends on its direction. Amplifying every channel by the same factor increases the first and leaves the second unchanged. Section 6 exhibits this concretely: with the Fisher score held equal to seven digits, the loss still spans a factor of 9.4.

5 The best achievable loss as a fractional control problem

Theorem 4.1 bounds the loss along a given trajectory. But the input steers the trajectory as well as the misalignment on it, so the best achievable loss requires optimising over both at once. This changes the structure of the problem.

The objective is a time average over a lifetime that the input itself determines:

$$\bar{\ell}^*(q) = \min_{u(\cdot) \in \mathcal{U}} \frac{\int_0^{T_f} \ell(\beta(x_t, u_t)) dt}{T_f}. \quad (24)$$

Both the integrand and the upper limit T_f depend on $u(\cdot)$. This makes (24) a fractional optimal control problem, meaning that the objective is a ratio rather than that the dynamics carry a fractional-order derivative. The fraction is not a technicality: it is the source of the trade-off, as follows.

Suppose T_f were fixed. Nothing in (2) penalises the input directly, so the answer would be trivial: find whichever admissible u minimises β and stay there. Because T_f is itself controlled, the denominator of (24) rules that solution out. The operating point that aligns

the indicator is in general not the operating point that prolongs life, so aligning now shortens the interval over which the alignment is averaged. The optimum must trade the two, and Section 6.5 measures how steep that trade is on a concrete plant. A fixed-horizon formulation represents none of it, which is why one was tried and discarded (Section 6.5).

5.1 A converse bound uniform over trajectories

Before computing anything we record a certificate that removes the conditioning. Theorem 4.1 bounds the loss along a trajectory the input has already chosen; the following proposition bounds it over all of them at once, and it is what would make the number computed below a limit rather than an observation, were a function satisfying it exhibited; none is, for reasons given below.

Proposition 5.1 (verification certificate). *Fix an indicator q and a constant $\lambda \geq 0$. Suppose there is a continuously differentiable V on a neighbourhood of the admissible state region such that*

(i) $V(x) \leq 0$ for every x on the failure surface $\{\max_i x_i = 1\}$, and

(ii) writing $D_i(x, u) = \varphi_i(u)h_i(x_i)$ for the drift of (2), for every admissible state x and every $u \in \mathcal{U}$,

$$\ell(\beta(x, u)) - \lambda + \nabla V(x)^\top D(x, u) \geq 0. \quad (25)$$

Then every admissible input satisfies $\bar{\ell} \geq \lambda + V(x_0)/T_f$. In particular, if $V(x_0) \geq 0$ then $\bar{\ell}^(q) \geq \lambda$, uniformly over all admissible inputs and hence over all trajectories.*

Proof. Let $u(\cdot)$ be any admissible input and x_t its trajectory. Since V is C^1 and $t \mapsto x_t$ is absolutely continuous, $t \mapsto V(x_t)$ is absolutely continuous and the chain rule gives $\frac{d}{dt}V(x_t) = \nabla V(x_t)^\top \dot{x}_t = \nabla V(x_t)^\top D(x_t, u_t)$ for almost every t . Applying (25) at (x_t, u_t) gives $\ell(\beta(x_t, u_t)) - \lambda + \frac{d}{dt}V(x_t) \geq 0$ a.e., and integrating over $[0, T_f]$,

$$\int_0^{T_f} (\ell - \lambda) dt \geq V(x_0) - V(x_{T_f}) \geq V(x_0),$$

the last step by (i), since x_{T_f} lies on the failure surface. Dividing by $T_f > 0$ gives the claim. $\square$

Condition (25) constrains the plant data alone: it quantifies over every state and every admissible input, and mentions no trajectory. Any V one can exhibit therefore certifies a limit, and the search for a converse becomes the search for a verification function.

The C^1 hypothesis is not a formality, and it is exactly what the natural candidates lack: the value function of the Dinkelbach problem below is only Lipschitz and violates (25) pointwise, while affine functions are smooth but certify only a vacuous pointwise bound. Supplementary Section S3 records both attempts and what they imply, that a certificate must be smooth and must have a gradient that varies, which points at a smooth under-approximation of the value function by sum-of-squares programming or interval arithmetic. We do not claim a certified converse here, and the optimum computed below remains a computation; what Proposition 5.1 buys is the replacement of an open problem without shape by a concrete object with a checkable property.

Problems of the form (24) are handled by Dinkelbach’s transform [27]: setting

$$F(\lambda) = \min_{u(\cdot) \in \mathcal{U}} \int_0^{T_f} (\ell - \lambda) dt, \quad (26)$$

the map F is decreasing in λ and vanishes exactly at the optimal ratio, so λ^* is found by bisection on $F(\lambda) = 0$. The inner problem in (26) has free terminal time and an absorbing set, hence a stationary Hamilton–Jacobi–Bellman equation on the state alone, with no time argument. Because $\dot{x} > 0$ the state advances monotonically, and a semi-Lagrangian scheme [28] converges by value iteration. One implementation detail matters: the scheme must use a fixed spatial step rather than a fixed time step, because degradation speeds vary by orders of magnitude across $\mathcal{U}$ and a fixed time step is unusable at one end of the envelope or the other.

6 Numerical study

Sections 3–5 are analytical, and three of their claims are stated sharply enough to be tested rather than illustrated: that the loss is invisible to a Fisher criterion, that the counting rule $m \geq d - 1$ binds in both directions on a generic plant, and that the best achievable loss is a finite number one can compute. This section instantiates the model on a synthetic three-mechanism plant on which every quantity is available in closed form, checks those three claims against it, measures what steering is worth, and asks which quantity, available in advance, says whether it will repay the effort. Section 7 then reports what measured fleets can and cannot establish about the premises.

Table 2: Loss of the fixed indicator at a Fisher score held equal. The exposure factor δ , defined in the text, holds $\delta^2 \|\mathbf{d}\|^2$ fixed and leaves $\hat{\mathbf{d}}(\tau)$ unaffected. Losses use $\gamma = \|\mathbf{d}\|^2 \text{Var}[R]$ with $\text{Var}[R] = 10^3$.

u	δ	$\delta^2 \ \mathbf{d}\ ^2$	rms β (deg)	$\bar{\ell}$ (nat)
0.0	0.235	0.4307565	6.23	0.0215
1.0	0.628	0.4307565	4.53	0.0123
1.5	1.000	0.4307565	6.22	0.0208
2.0	1.564	0.4307565	7.53	0.0379
3.0	3.769	0.4307565	6.57	0.1161

6.1 The test plant

We take $d = 3$ mechanisms with Arrhenius factors $\varphi_i(u) = e^{\theta_i - E_i u}$, in which the input u is the reciprocal temperature $1/(kT)$, so that this is the Arrhenius form of Section 2 with e^{θ_i} the pre-exponential constant A_i . We set $\theta = (0, 0.3, -0.2)$ and $E = (0.6, 1.1, 0.8)$, and shapes $h_i(x) = 1 + c_i x$ with $c = (-0.5, 2.0, 0.8)$, giving one decelerating and two accelerating mechanisms. There is a single input $u \in [0, 3]$, so $m = 1 < d - 1 = 2$, and Theorem 4.1 puts zero loss out of reach unless this plant is one of the exceptions of Remark 4.5; Section 6.3 computes the floor and finds that it is not. Noise is unit and the prior variance is $\text{Var}[R] = 10^3$, so the weight $w = \gamma/(1 + \gamma)$ of Lemma 2.2 is evaluated at the plant’s own $\gamma = \|\mathbf{d}\|^2 \text{Var}[R]$ at each state rather than held fixed; γ varies along a trajectory because $\|\mathbf{d}\|$ does. Every trajectory starts from $x_i(0) = 5 \times 10^{-3}$ and runs to the condemnation limit $x_f = 0.9$ of (3). Raising x_f to unity leaves every qualitative conclusion below intact and moves the constants: the peak floor of Section 6.3 rises from 4.52° to 5.41° , and the gain of Section 6.5 falls from 68 to 58. No constant reported below should be read across thresholds.

Verification of the identities The analytical identities were checked numerically before use, and the residuals are reported in Supplementary Section S4.

6.2 A Fisher level set of varying loss

Proposition 4.6 is easy to demonstrate on this plant. For a constant input, a common gain on $\mathbf{d}$ moves $\|\mathbf{d}\|$ while leaving $\hat{\mathbf{d}}(\tau)$, and hence β , untouched: the Fisher score and the loss can be varied independently. Write $\delta(u)$ for the gain that equalises the score across inputs, fixed by $\delta(u)^2 \|\mathbf{d}\|^2(u) = F$ with F its value at $u = 1.5$. It is realised by rescaling how long the component runs under load, which is why it does not rotate $\hat{\mathbf{d}}$. By construction $\delta = 1$ at $u = 1.5$, and $\delta > 1$ wherever the raw score is smaller, so it is a multiplier on exposure rather than a fraction of it. Table 2 does exactly that, holding the Fisher score equal to seven digits by tuning the exposure factor δ . The loss still spans a factor 9.4.

The two objectives also disagree on which input to prefer: the raw score, before δ equalises it, is largest at




Figure 4: (a) The information-share curve p_τ on the 2-simplex (circle: start, square: failure), the reachable sets $\mathcal{D}(x_\tau)$ at $\tau = 0.05, 0.45$ and 0.9 , and the optimal indicator q^* . No point lies on all three, so β cannot be held at zero (Theorem 4.1). (b) What steering buys, against ρ , the drift of the reachable arc over its length.




Figure 5: (a) With $m = 1$ the floor against q^* is largest at birth and dips below the sampling grid near mid-life. (b) With $m = 2$ the mid-life direction is held exactly; plotted is the operating point this requires, which never leaves $\mathcal{U}$.

6.3 The reachability floor with a single input

With $m = 1 < d - 1$ no direction lies on every reachable set at once. Along the closed loop at q^* the floor is largest at the start of life, where $\beta_{\min} = 4.52^\circ$; it falls through mid-life to below the 0.01° resolution of our sampling, where the indicator is momentarily reachable, and rises again towards failure (Fig. 5(a)). Zero average loss would need the indicator reachable at every instant and not at one, so it is unattainable however the input is steered, in agreement with Theorem 4.1. Fig. 4(a) shows the geometry: the reachable set sweeps across q^* as the plant ages, so the mid-life arc passes close to the indicator while the arcs at $\tau = 0.05$ and $\tau = 0.9$ stand off it.

6.4 The counting rule, tested in both directions

Theorem 4.1 is an equivalence, and the paragraph above tests only half of it: with too few inputs the floor stays positive. The other half, that supplying the inputs the counting rule asks for makes the floor vanish, is the more useful claim for a practitioner deciding how many operating variables to instrument, so we test it directly.

Give the same three-mechanism plant a second independent input by writing $\varphi_i(u) = \exp(\theta_i - E_i u_1 - G_i u_2)$ with $G = (0.9, 0.3, 1.2)$, so that $m = 2 = d - 1$ and each input acts on the mechanisms in a different pattern. Then $\log d_i$ is affine in (u_1, u_2) , and matching a target direction fixes the $d - 1 = 2$ differences of logarithms: two linear equations in two unknowns, with a coefficient matrix of determinant 0.27 here, hence one unique solution at every state.

Taking as target the direction the plant occupies at

mid-life under a central operating point, that solution lies inside the admissible box $[0, 3]^2$ at every one of 600 sampled states covering the whole life, and the residual misalignment there is 8.5×10^{-7} degrees, that is, zero to numerical precision; the schedule it requires is drawn in Fig. 5(b). An independent 90×90 grid search over the box returns a floor of at most 0.156° , which is the resolution of the grid rather than a property of the plant.

The contrast with the single-input case is the content of the theorem: the same mechanisms, the same shapes, the same envelope on each input, and the peak floor moves from 4.52° to zero at every sampled state on the addition of one operating variable. On this plant the counting rule is therefore not a conservative sufficient condition but a sharp one, and within the adopted model it fixes how many independent control degrees of freedom the plant must offer. How those degrees of freedom correspond to physical actuators, and how the mechanisms are individuated in the first place, are modelling questions this paper does not settle. Proposition 4.4 says which plants escape it, and Remark 4.5 says why a designer should not count on being one of them.

6.5 The value of steering

Minimising over q gives $q^* = (0.466, 0.771, 0.435)$, at which the adaptive policy attains 1.68×10^{-4} against 1.15×10^{-2} under the best constant input, a factor of 68. The ratio depends on how finely $\mathcal{U}$ is discretised for the policy, rising to 72 at 161 points; we report it on the 41-point grid used throughout.

The factor is a ratio of information losses, not of information, and it is not free in reliability terms: the policy that attains it runs the component differently, and Section 6.6 finds the two closed loops separated by a factor of 3.4 in life. Aligning the indicator and prolonging the component are competing objectives, which is what makes the horizon in (24) a controlled quantity rather than a constant. That factor is the value of steering at all, not the value of steering optimally: a myopic policy that attends only to the numerator of (24) sits 2.8% above the optimum at q^* .

The fraction in (24) is doing visible work at this op-

erating point. The closed loop at q^* and the closed loop at a balanced indicator, which calls for far less steering, differ by a factor 3.4 in lifetime: what alignment buys in the numerator is paid for in the denominator. This is the tension anticipated in Section 5, and it is why the horizon cannot be fixed in advance.

6.6 The computed fractional optimum

Dinkelbach bisection with the control grid matched to the policy gives, over $N \in \{25, 29, 35\}$, extrapolants agreeing to 1.2% and $\lambda_\infty^* = 1.63 \times 10^{-4}$, below the realisable 1.68×10^{-4} as it must be.

We quote this to two digits only, for the following reason: the raw values are twice the extrapolant, so the first-order correction carries half the answer, and both quantities still move by 4% between control grids of 41 and 81 points. Three checks guard the number: F must decrease and cross zero, λ^* must not exceed a realisable policy, and the grid sequence must extrapolate consistently. A fixed-horizon formulation failed the second and third and was discarded, as did a first attempt with mismatched control grids. A converse bound uniform over trajectories, which would replace this computation by a theorem, remains open.

6.7 Authority, drift, and how far they screen

For a practitioner the first question is not how to steer but whether to bother. The reachable set $\mathcal{D}(x)$ is an arc, and it offers two natural answers: its mean length, which is the actuation authority available to chase the informative direction, and the drift of its midpoint over the life, which is how far that direction wanders. Write ρ for the drift divided by the length, which weighs the two against each other.

Table 3 varies the spread of the activation energies, which by (20) is what sets that authority. As the arc collapses, the attainable loss L^* rises by a factor 25 and the value of control falls from 68 to 2.7: with weak authority there is little left to steer. Fig. 4(b) plots the two against ρ . In this family the drift barely moves, staying between 24° and 30° while the arc spans nearly two decades, so ρ is the reciprocal of the arc in all but name and the table cannot say which of the two matters. Both can be computed from a nominal degradation model before any indicator is designed.

Three plants differing only in their activation energies are not a screening law, so the index was tested on 299 further plants drawn at random: two to five mechanisms, one or two independently adjustable inputs, and the pre-exponentials, activation energies, shape coefficients, input envelope, initial state and signal-to-noise ratio varied together. The rank correlation between ρ and the gain is -0.36 , with a bootstrap 95% interval from -0.46 to -0.24 , and it keeps its sign inside every

Table 3: Authority and drift in a one-parameter family. The activation energies E set the actuation authority; c , the noise and $\mathcal{U}$ are held fixed. L^* is the loss under adaptive control, and the gain its ratio to the best constant input.

E	arc	ρ	L^*	gain
(0.60, 1.10, 0.80)	33.9°	0.88	1.7×10^{-4}	$68\times$
(0.80, 0.82, 0.79)	1.93°	12.5	3.6×10^{-3}	$2.9\times$
(0.80, 0.805, 0.798)	0.45°	53.1	4.2×10^{-3}	$2.7\times$

stratum of the mechanism count and of the input count, so it is not an artefact of either.

Read as a classifier of whether steering repays a factor of five, which 32% of the plants do, ρ has an area under the receiver-operating curve of 0.73, with interval 0.67 to 0.79. The reachable arc alone scores 0.73 as well, and the drift alone 0.53, with interval 0.46 to 0.60, which includes the 0.5 of a coin toss.

Equal areas do not make the two equally good, and the question is settled by a model that contains both as special cases: a logistic regression on the logarithms of arc and drift as separate terms, of which the arc alone and the ratio ρ are each a constrained case. The drift adds nothing to the arc as a main effect ($p = 0.47$), and constraining the two to enter as a ratio fits significantly worse than the arc alone ($p = 7.7 \times 10^{-5}$; Akaike information 351 against 336). Both results hold in each input count separately. The drift does carry information, but through the arc ($p = 3.0 \times 10^{-4}$ for their interaction), and in opposite directions at the two ends. Among plants of low authority the share repaying fivefold falls from 0.18 to 0.08 as the drift rises from its lowest to its highest third; among plants of high authority it rises from 0.48 to 0.61. Drift that outruns the authority defeats steering, and drift that the authority can follow is what gives steering something to do. The ratio encodes only the first of these. It is also why the lowest tenth of ρ repays fivefold only 43% of the time: there the drift is small, a median of 2.7° against 14° over the sample.

Out of sample the ranking agrees. Over repeated tenfold cross-validation the arc and ρ each rank unseen plants with an area of 0.72, and a model adding the drift and its interaction with the arc reaches 0.74, higher in every repeat. The quantity to screen with is therefore the arc, the authority itself, refined by the drift only through that interaction; ρ remains a convenient description of the family of Table 3, where the drift is nearly fixed. The first row of that table is not typical of its own regime: among sampled plants within a factor 1.4 of its ρ , the median gain is 5.2 rather than 68.

6.8 What the loss buys a downstream estimator

Corollary 2.4 prices the loss in the error of the estimate the model itself prescribes. A practitioner does not use that estimate, so the ordering is worth checking against a

procedure that uses none of the construction: sample the indicator over the last 40% of life, match the observed history against the indicator's own reference curve, and read remaining life off the best-fitting time offset. That estimator sees a scalar signal and nothing else — not ℓ , not γ , not the loss.

Over twelve indicators spanning $\bar{\ell}$ from 0.012 to 0.37, its median absolute error — one of the standard prognostic accuracy measures [29] — rises with the loss, at Spearman correlation 0.89, and the best-aligned indicator beats the worst by a factor 1.5. Twelve indicators is a small sample, so both are given with intervals obtained by resampling the 160 units behind each median: 0.34 to 0.90 for the correlation, and 1.40 to 2.03 for the factor. The ordering is not in doubt — a one-sided permutation test gives $p = 1.5 \times 10^{-4}$ — but its strength is, and the point estimate sits near the top of its own interval. The factor is far smaller than the 31-fold spread in $\bar{\ell}$, and it should be: the loss is a logarithmic information quantity, the error is a time, and an estimator with the whole history recovers much of what any one instant discards.

The same protocol on 120 further plants, drawn at random as in Section 6.7, gives correlations between 0.87 and 0.99, with the plant above among the weaker. That design tilts every indicator toward the same direction, however, so loss and error both grow with one parameter and their agreement is close to structural. Tilting instead by a random amount of up to 60° in a random direction weakens the median correlation to 0.74, with the interval excluding zero on 84% of plants. It also exposes what the experiment cannot show. On these plants the angle between an indicator and the mid-life informative direction ranks the indicators identically to the loss on 82% of plants, with neither consistently better on the rest, because the direction rotates little beside how the indicators differ. The experiment is therefore evidence that the ordering of the loss carries over to an estimator built from none of this construction; it is not evidence that the time-averaged loss improves on a mid-life snapshot of the misalignment, which would need a plant whose direction rotates by more than its candidate indicators differ.

7 What public data can say about the premises

The construction rests on three modelling choices, and each is empirically testable. They do not carry equal weight, and Fig. 6 separates what depends on which. Two are premises about the plant: that the informative direction rotates along the trajectory, and that the input moves it. The third is the separability of Lemma 4.3, which is what makes the operator's authority computable offline. Table 4 summarises what six public run-to-failure fleets currently settle. Supplementary Section S1 gives the protocol behind every row, and its Table S1 lists the parameters of the code that produced it.

Three readings qualify the table. The Severson channels are capacity fade, resistance growth and mean temperature, and mean temperature is a monitoring channel rather than a damage state — it is the operating variable itself in the Arrhenius form of Section 2. That row therefore reads the rotation of a direction formed from three measured series, which is what the premise asserts, and does not need (5)'s pairing of channels with mechanisms. On the four measured fleets the obstacle to testing actuation is the estimator and not the physics, so the premise is supported in simulation and untested on hardware. Both the Severson and the C-MAPSS rows estimate the direction by differentiating each channel along normalised life, which reads a time derivative as the sensitivity of Eq. (6) and so presumes the time-shift population of Section 2. And the C-MAPSS contrast is reported against a negative control — fake regimes cut from a single real one, where no contrast exists by construction — because a procedure that returned a positive contrast regardless would produce the same table entry.

The consequence is the scope of the results above. The geometry of Sections 3 and 4 needs only $\varphi_i > 0$ and $h_i > 0$, and is untouched. What separability buys, and what the rejection withdraws, is Proposition 4.4 and the offline route to $\mathcal{B}$: those describe a model class, and a plant must be shown to belong to it before they are used. The two-way layout above is the test, and it costs one pass over data a reliability programme already holds.

8 Discussion

Four practical readings follow from the results above. The first two hold whatever the plant is, because one is a test and the other an invariance. The last two assume the separable form of (2), which Section 7 reports as rejected on the only data able to test it, so they apply where that form does and not otherwise.

Test the separability assumption before relying on it. The two-way layout of Section 7 needs only a fleet whose operating point varies within a unit, and costs one pass over data a reliability programme already holds. If the interaction clears the measured noise floor, the offline route to the operator's authority is unavailable and the two model-based readings below do not apply.

No Fisher criterion can rank fixed indicators by the loss they incur. This is the second of the two unconditional readings: Proposition 4.6 is an invariance of the loss under a common gain and holds whatever the plant, so rules that maximise a scalarisation of the Fisher information cannot separate configurations whose losses differ by an order of magnitude. Where the two objectives do coincide, that is a property of the application, to be verified rather than assumed.

The binding constraint is actuation authority, not indicator sophistication. Theorem 4.1 prices the operating variables a lossless fixed indicator would need, and Proposition 4.4 sharpens the count to sufficient but not necessary, with exceptions that Remark 4.5 shows to be




Figure 6: What rests on what. The spine needs only the positivity of the degradation model. Only the dashed branch adds separability, which Section 7 reports as rejected on the best public turbofan data.

Table 4: What the six public fleets settle, and what they cannot. Only the two turbofan simulations vary the operating point within a unit, which is why the actuation premise is supported in simulation and untested on measured hardware.

fleet	premise	what is measured	reading
Severson lithium-ion [30]	rotation	arc of an empirically constructed informative direction, the fleet mean, over $\tau \in [0.15, 0.9]$, on 129 Severson lithium-ion cells	travels 30.7° of arc and ends 21° from its start (Fig. 2(c)); stays in the open positive orthant once the window reaches 301 cycles
XJTU-SY [31], PRONOSTIA [32]	actuation	—	the operating condition is fixed per unit, so condition and group of units cannot be separated four cells
NASA batteries [33]	actuation	—	by 7.9° (FD004) and 10.8° (FD002), positive for 436 of 439 engines; fake regimes give $ B - W \leq 0.5^\circ$
C-MAPSS [34]	actuation	between-regime minus within-regime angle, each engine serving as its own control	larger by a factor 1.65 over the operating-point effect on DS02 and 1.84 on DS01, under every clustering and staging tried: rejected
N-CMAPSS DS02, DS01 [35]	separability	interaction against both main effects, in a two-way layout inside engines	

a knife edge. Care in choosing the weights substitutes for actuation in neither case. What one controllable variable does buy is the weaker objective of Corollary 4.2: an installation that could never freeze the informative direction can still hold the misalignment constant, though not at zero. Authority bought this way is not free, and Section 6.6 prices it in life.

The decision to steer precedes the design of the indicator. The width of the reachable arc requires only a nominal degradation model and the operating envelope; and Table 3 shows the value of steering collapsing from $68\times$ to under $3\times$ as it shortens. It screens rather than decides, and across random plants the drift of the informative direction refines it only in combination: drift the authority can follow gives steering something to do, and drift that outruns it defeats steering (Section 6.7).

Limitations Four limitations qualify the readings above. The separable form of (2) is not only an idealisation but one the data reject: Section 7 finds the interaction between operating point and state larger than either main effect on two N-CMAPSS datasets, so computing the operator’s authority offline is available only where that form holds. The optimum of Section 5 is computed rather than certified, to two digits. The actuation premise is supported by simulated rather than measured data, for the reasons set out in Section 7. And the population is taken to slide along one degradation path, which is what makes remaining life a coordinate on it; under dispersion in the coefficients of (2) the sensitivity that Remark 2.1 measures stands some 45° from (6), so the construction describes fleets that vary in when they fail rather than in how they degrade.

9 Conclusion

A fixed health indicator discards information, and this paper has treated what it discards as a quantity the operating point controls. Indicator design then becomes a reachability question. The family of reachable direction sets bounds what the operating point can achieve, and their width, available before any indicator is fixed, indicates whether substantial benefit is plausible, and across random plants of this model the drift of the informative direction refines that indication only in combination with it. Under the proposed model, actuation authority, rather than sophistication of the indicator, governs whether a fixed compression can come close to lossless. Whether that authority is present is an empirical question rather than a modelling assumption, and on 439 simulated turbofan records the evidence supports it. Establishing the same on measured fleets awaits records in which the operating point varies enough to show the effect.

Section 7 rejects that separable form on the highest-fidelity public turbofan data, so the geometry of Section 4.1, and the offline route it opens, describe a model class rather than measured machinery. What replaces that structure when the stress model is not separable, and whether a weaker one still yields a computable reachability criterion, remains open.

Two further questions remain open, the first of them now in a more specific form. Proposition 5.1 reduces the search for a converse bound uniform over trajectories to the exhibition of a single continuously differentiable verification function; the value function is the natural candidate and fails the differential test for a structural reason, so what is required is a smooth underapproximation of it, by sum-of-squares programming or by interval arithmetic with an explicit Lipschitz constant. The other is unchanged: the computed optimum to better than two digits, which requires a scheme of higher order than the semi-Lagrangian one used here.

Data availability and reproducibility

The six fleets named in Section 7 are publicly available. One script recomputes every reported number from a single configuration and compares it with the value written in this manuscript and in its supplementary material, covering the tables and the quantities reported alongside them; all of its checks pass, and the output of the run that verified this version is included. The random-plant sweep of Section 6.7 takes hours, so it and the downstream repeat of Section 6.8 are recorded by two further scripts, and the checks read the records. Three groups of numbers lie outside the checks: the residuals of the identity checks in Supplementary Section S4, the barycentre comparison in Section 3.3, and the diagnostics of the two candidate certificates in Supplementary Section S3. The plots are produced by three further scripts. All of these are available

at <https://github.com/HuyHoangLe0201/health-indicator-information-loss>, and are also provided with this article as a supplementary archive, together with the Supplementary Material.

Supporting information

Additional supporting information is found in the online version of this article.

Supplementary Material (PDF). Supplementary Sections S1 to S4: the protocols behind every row of Table 4 with the parameters of the code that produced them, the proofs of Propositions 2.3 and 3.3, the two candidate certificates for Proposition 5.1 and why neither qualifies, and the numerical verification of the identities and of the exceptional family of Remark 4.5.

Supplementary archive (ZIP). The Supplementary Material and a copy of the repository named in the Data availability section: the script that recomputes and checks every reported number, the records of the two sweeps it reads, the output of the run that verified this version, and the three figure scripts.

Conflict of interest

The authors declare that they have no known competing financial interests or personal relationships that could have appeared to influence the work reported in this paper.

Funding

This research did not receive any specific grant from funding agencies in the public, commercial, or not-for-profit sectors.

References

- [1] Yaguo Lei, Naipeng Li, Liang Guo, Ningbo Li, Tao Yan, and Jing Lin. Machinery health prognostics: A systematic review from data acquisition to RUL prediction. *Mech. Syst. Signal Process.*, 104:799–834, 2018. doi: 10.1016/j.ymssp.2017.11.016.
- [2] Jamie Coble and J. Wesley Hines. Identifying optimal prognostic parameters from data: A genetic algorithms approach. In *Annu. Conf. PHM Soc.*, 2009.
- [3] Huy Hoang Le and Kim-Anh Nguyen. Information-theoretic limits of health indicator construction for remaining useful life prediction. *Eng. Sci. Technol. Int. J.*, 81:102485, 2026. doi: 10.1016/j.jestch.2026.102485.

- [4] Huy Hoang Le and Kim-Anh Nguyen. Generalization bounds and sample complexity for remaining useful life prediction from complete degradation trajectories. *Meas. Sci. Technol.*, 37:226203, 2026. doi: 10.1088/1361-6501/ae7109.
- [5] Kaibo Liu, Nagi Z. Gebraeel, and Jianjun Shi. A data-level fusion model for developing composite health indices for degradation modeling and prognostic analysis. *IEEE Trans. Autom. Sci. Eng.*, 10(3):652–664, 2013. doi: 10.1109/TASE.2013.2250282.
- [6] Changyue Song and Kaibo Liu. Statistical degradation modeling and prognostics of multiple sensor signals via data fusion: A composite health index approach. *IIEE Trans.*, 50(10):853–867, 2018. doi: 10.1080/24725854.2018.1440673.
- [7] Liang Guo, Naipeng Li, Feng Jia, Yaguo Lei, and Jing Lin. A recurrent neural network based health indicator for remaining useful life prediction of bearings. *Neurocomputing*, 240:98–109, 2017. doi: 10.1016/j.neucom.2017.02.045.
- [8] Kamran Javed, Rafael Gouriveau, and Noureddine Zerhouni. State of the art and taxonomy of prognostics approaches, trends of prognostics applications and open issues towards maturity at different technology readiness levels. *Mech. Syst. Signal Process.*, 94:214–236, 2017. doi: 10.1016/j.ymssp.2017.01.050.
- [9] Zhi-Sheng Ye and Min Xie. Stochastic modelling and analysis of degradation for highly reliable products. *Appl. Stoch. Models Bus. Ind.*, 31(1):16–32, 2015. doi: 10.1002/asmb.2063.
- [10] J. M. van Noortwijk. A survey of the application of gamma processes in maintenance. *Reliab. Eng. Syst. Saf.*, 94(1):2–21, 2009. doi: 10.1016/j.res.2007.03.019.
- [11] Xiao-Sheng Si, Wenbin Wang, Chang-Hua Hu, and Dong-Hua Zhou. Remaining useful life estimation — a review on the statistical data driven approaches. *Eur. J. Oper. Res.*, 213(1):1–14, 2011. doi: 10.1016/j.ejor.2010.11.018.
- [12] Zhengxin Zhang, Xiaosheng Si, Changhua Hu, and Yaguo Lei. Degradation data analysis and remaining useful life estimation: a review on Wiener-process-based methods. *Eur. J. Oper. Res.*, 271(3):775–796, 2018. doi: 10.1016/j.ejor.2018.02.033.
- [13] Jay Lee, Fangji Wu, Wenyu Zhao, Masoud Ghafari, Linxia Liao, and David Siegel. Prognostics and health management design for rotary machinery systems — reviews, methodology and applications. *Mech. Syst. Signal Process.*, 42(1–2):314–334, 2014. doi: 10.1016/j.ymssp.2013.06.004.
- [14] Fatemeh Karimi Pour, Didier Theilliol, Vicenç Puig, and Gabriela Cembrano. Health-aware control design based on remaining useful life estimation for autonomous racing vehicle. *ISA Trans.*, 113:196–209, 2021. doi: 10.1016/j.isatra.2020.03.032.
- [15] M. S. Félix, J. J. Martinez, and C. Bérenguer. Remaining useful life (RUL) control of controlled systems under degradation. *Int. J. Robust Nonlinear Control*, 35(18):8095–8116, 2025. doi: 10.1002/rnc.70105.
- [16] A. A. Feldbaum. Dual control theory I. *Autom. Remote Control*, 21(9):874–880, 1960.
- [17] Christopher Grebe, Emmett Wise, and Jonathan Kelly. Observability-aware trajectory optimization: Theory, viability, and state of the art. In *IEEE Int. Conf. Multisensor Fusion Integr. (MFI)*, pages 1–8, 2021. doi: 10.1109/mfi52462.2021.9591177.
- [18] Giles Hooker, Kevin K. Lin, and Bruce Rogers. Control theory and experimental design in diffusion processes. *SIAM/ASA J. Uncertain. Quantif.*, 3(1):234–264, 2015. doi: 10.1137/140962280.
- [19] Heonsang Lim and Bong-Jin Yum. Optimal design of accelerated degradation tests based on Wiener process models. *J. Appl. Stat.*, 38(2):309–325, 2011. doi: 10.1080/02664760903406488.
- [20] Shun-ichi Amari. *Information Geometry and Its Applications*. Springer, 2016. doi: 10.1007/978-4-431-55978-8.
- [21] Anil Bhattacharyya. On a measure of divergence between two statistical populations defined by their probability distributions. *Bull. Calcutta Math. Soc.*, 35:99–109, 1943.
- [22] Josef Hofbauer and Karl Sigmund. *Evolutionary Games and Population Dynamics*. Cambridge Univ. Press, 1998. doi: 10.1017/cbo9781139173179.
- [23] William H. Sandholm. *Population Games and Evolutionary Dynamics*. MIT Press, 2010.
- [24] William Q. Meeker and Luis A. Escobar. *Statistical Methods for Reliability Data*. Wiley, 1998.
- [25] Thomas M. Cover and Joy A. Thomas. *Elements of Information Theory*. Wiley, 2nd edition, 2006. doi: 10.1002/047174882x.
- [26] Erkki Oja. A simplified neuron model as a principal component analyzer. *J. Math. Biol.*, 15:267–273, 1982. doi: 10.1007/bf00275687.
- [27] Werner Dinkelbach. On nonlinear fractional programming. *Manage. Sci.*, 13(7):492–498, 1967. doi: 10.1287/mnsc.13.7.492.

- [28] Maurizio Falcone and Roberto Ferretti. *Semi-Lagrangian Approximation Schemes for Linear and Hamilton–Jacobi Equations*. SIAM, 2014. doi: 10.1137/1.9781611973051.
- [29] Abhinav Saxena, Jose Celaya, Edward Balaban, Kai Goebel, Bhaskar Saha, Sankalita Saha, and Mark Schwabacher. Metrics for evaluating performance of prognostic techniques. In *International Conference on Prognostics and Health Management*, pages 1–17, 2008. doi: 10.1109/phm.2008.4711436.
- [30] Kristen A. Severson, Peter M. Attia, Norman Jin, Nicholas Perkins, Benben Jiang, Zi Yang, Michael H. Chen, Muratahan Aykol, Patrick K. Herring, Dimitrios Fraggidakis, Martin Z. Bazant, Stephen J. Harris, William C. Chueh, and Richard D. Braatz. Data-driven prediction of battery cycle life before capacity degradation. *Nat. Energy*, 4:383–391, 2019. doi: 10.1038/s41560-019-0356-8.
- [31] Biao Wang, Yaguo Lei, Naipeng Li, and Ningbo Li. A hybrid prognostics approach for estimating remaining useful life of rolling element bearings. *IEEE Trans. Rel.*, 69(1):401–412, 2020. doi: 10.1109/tr.2018.2882682.
- [32] Patrick Nectoux, Rafael Gouriveau, Kamal Medjaher, Emmanuel Ramasso, Brigitte Chebel-Morello, Nouredine Zerhouni, and Christophe Varnier. PRONOSTIA: An experimental platform for bearings accelerated degradation tests. In *IEEE Int. Conf. Prognostics Health Manage. (PHM)*, pages 1–8, Denver, CO, USA, 2012.
- [33] Bhaskar Saha and Kai Goebel. Battery data set. NASA Ames Prognostics Data Repository, 2007.
- [34] Abhinav Saxena, Kai Goebel, Don Simon, and Neil Eklund. Damage propagation modeling for aircraft engine run-to-failure simulation. In *Int. Conf. Prognostics Health Manage. (PHM)*, pages 1–9, Denver, CO, USA, 2008. doi: 10.1109/phm.2008.4711414.
- [35] Manuel Arias Chao, Chetan Kulkarni, Kai Goebel, and Olga Fink. Aircraft engine run-to-failure dataset under real flight conditions for prognostics and diagnostics. *Data*, 6(1):5, 2021. doi: 10.3390/data6010005.