

Supplementary material for: Information loss of fixed linear health indicators as a control problem: simplex geometry and reachability limits

Huy Hoang Le¹ and Kim-Anh Nguyen²

¹R&D Department, PowerMore Ltd., Da Nang 550000, Viet Nam

²Faculty of Electrical Engineering, University of Science and Technology, The University of Danang, Da Nang 550000, Viet Nam

Abstract

Section S1 gives the protocol behind every row of Table 4 of the main text, with the parameters of the code that produced it. Sections S2 to S4 hold material moved out of the main text: two proofs, the two candidate certificates for Proposition 5.1 and why neither qualifies, and the numerical verification of the identities and of the exceptional family of Remark 4.5. Equation, section and proposition numbers without an S refer to the main text.

S1 Protocols for the public-fleet tests

Every number in Table 4 is recomputed by `verify_numbers.py` in the repository named in the main text, and the parameters below are the ones that script passes; where a test was also run at other settings to check stability, both are given. All four protocols read an empirical informative direction as a time derivative of the monitoring channels along normalised life, and so rely on the time-shift population of Section 2, under which that derivative is the sensitivity of Eq. (6).

S1.1 The rotation test on the Severson cells

Channels are capacity at discharge, internal resistance and mean cell temperature. A cell is used if its life is recorded, it has at least 400 cycles, and all three channels are finite throughout; 129 cells qualify. Life is normalised per cell to $\tau \in [0, 1]$ by cycle index. Each channel is differentiated with a Savitzky–Golay filter of polynomial order 3 and window $\min(301, 2\lfloor n/4 \rfloor + 1)$ cycles for a cell of n cycles, and the derivative is divided by the standard deviation of that channel’s residual about the same filter, which whitens the channel by its own measured noise. The whitened derivative is interpolated onto 40 points of $\tau \in [0.15, 0.90]$. One sign per channel, the median sign over every cell and grid point, orients all cells alike; each cell’s vector is then normalised at every grid point, and the fleet mean, renormalised, is the empirical informative direction. Its arc is the sum of the angles between consecutive grid points, and its end-to-end angle is the angle between the first and the last. Positivity is the fraction of grid points at which every component of the fleet mean is positive; it is one at the window used.

S1.2 Why three measured fleets cannot test actuation

Testing actuation needs the operating point to change within a unit while everything else is held. In XJTU-SY each of three operating conditions is applied to its own five bearings, and in PRONOSTIA each of three conditions to its own bearings, so a difference between conditions cannot be told apart from a difference between groups of units. The NASA battery fleet has four cells. These fleets are therefore listed with the reason and no statistic.

S1.3 The actuation test on C-MAPSS

The training files of FD004 and FD002 are used. Life is normalised per engine by its final cycle. Operating regimes are labelled by standardising the three operating settings over the file and rounding to one decimal; each file has 6 regimes. A sensor is kept if its standard deviation within a (engine, regime) group has a positive median, and each kept sensor is standardised within each regime, so that a regime's level does not enter the direction.

Within one engine, each regime's cycles are resampled to a common count m , the smallest over that engine's regimes, and an engine enters only if $m \geq 2(w + 3)$. Each resampled series is split into its even and its odd cycles, two half-samples of the same regime. On each half the direction is estimated as for the Severson cells, with polynomial order 2 and window w , and interpolated onto 10 points of $\tau \in [0.30, 0.85]$. The angle between two directions is the mean over grid points of the pointwise angle. W is the mean, over the engine's regimes, of the angle between the two halves of the same regime, and B the mean, over pairs of different regimes, of the angle between a half of one and a half of the other. $B - W$ is positive when regimes differ by more than half-samples of one regime differ from each other, which is what an operating point that moves the direction predicts. The table reports the mean of $B - W$ over engines, and the number of engines for which it is positive, at $w = 7$; at $w = 9$ and $w = 11$ every value keeps its sign and exceeds the largest control bias.

The negative control uses the same machinery on fake regimes. Within one regime of one engine with at least $4(w + 3)$ cycles, the four interleaved subsequences of every fourth cycle stand in for four regimes; W compares the first with the second and B the first with the fourth. No contrast exists by construction. Over all three windows and both files the mean of $B - W$ stays within 0.5° of zero, and a paired t test of B against W does not reject equality at any window, the smallest p being 0.35.

S1.4 The separability test on N-CMAPSS

The development sets of DS02, units 2, 5, 10, 16, 18 and 20, and of DS01, units 1 to 6, are used. The operating point is held fixed by clustering flight points, by k -means with 6 clusters, on standardised altitude, Mach number and throttle-resolver angle; life is divided into 2 stages of equal width in cycles. A (stage, cluster) cell enters only with at least 400 points.

Within a cell each sensor is regressed on standardised remaining life, and the slope is divided by the standard deviation of that sensor's residual, so that the vector of slopes is the informative direction at a fixed operating point rather than a residual after standardisation. The vector is normalised, and the logarithms of the absolute values of its components, centred on their mean, are what the test compares. The logarithm is the point: under the separable form of Lemma 4.3 each component is a product of an operating-point factor and a state factor, so in logarithms the two act additively, and separability is the absence of interaction.

Each cell is estimated twice, from its even and from its odd points. The mean of the two is the signal; their half-difference measures noise, and applying the same decomposition to it

Table S1: The parameters of each test as run.

test	parameters
Severson rotation	129 cells with at least 400 cycles; Savitzky–Golay order 3, window 301 cycles or $2\lfloor n/4 \rfloor + 1$ if smaller; 40 points on $\tau \in [0.15, 0.90]$
C-MAPSS actuation	FD004 and FD002 training sets; 6 regimes each; order 2, window $w = 7$, stability at $w = 9, 11$; 10 points on $\tau \in [0.30, 0.85]$; engine enters if $m \geq 2(w + 3)$; control needs $4(w + 3)$ cycles in one regime
N-CMAPSS separability	DS02 units 2, 5, 10, 16, 18, 20 and DS01 units 1–6; 6 clusters, 2 stages, at least 400 points per cell; sweep over seven (cluster, stage) settings with at least 300 points

gives each effect its own noise floor. The decomposition is the ordinary two-way one over the cells, into a grand mean, a stage effect, an operating-point effect and an interaction, with sums of squares taken over cells and sensors and pooled over units. With 6 clusters and 2 stages the operating-point effect and the interaction each have 5 degrees of freedom per sensor and the stage effect 1, so the reported ratio of interaction to operating-point effect compares quantities with equal degrees of freedom. That ratio is 1.65 for DS02 and 1.84 for DS01. Against its own noise floor the interaction is 566 times larger on DS02 and 110 times on DS01, so it is not noise. The test was repeated at 3, 4, 6, 8 and 10 clusters with 2 stages, at 6 clusters with 3 stages and at 4 clusters with 4 stages, with at least 300 points per cell; the interaction exceeds the operating-point effect in every configuration, by a factor of at least 1.42.

S2 Proofs of Propositions 2.3 and 3.3

Proof of Proposition 2.3. In whitened coordinates the linearised record is $\tilde{y}_k = \mathbf{d}_k R + \epsilon_k$ with $\epsilon_k \sim N(0, I_d)$ independent, so the Fisher information it carries about R is $\sum_k \mathbf{d}_k^\top \mathbf{d}_k = \sum_k \|\mathbf{d}_k\|^2$ and an oracle reading the whole record obtains $\frac{1}{2} \log(1 + G)$ nats. The fixed indicator returns the scalars $z_k = \mathbf{v}^\top \tilde{y}_k = (\mathbf{v}^\top \mathbf{d}_k) R + \mathbf{v}^\top \epsilon_k$, whose noise terms $\mathbf{v}^\top \epsilon_k$ are standard normal and independent because $\|\mathbf{v}\| = 1$, so the compressed record carries Fisher information $\sum_k (\mathbf{v}^\top \mathbf{d}_k)^2$. Since $\mathbf{v}^\top \mathbf{d}_k = \|\mathbf{d}_k\| \cos \psi_k$,

$$\text{Var}[R] \sum_k (\mathbf{v}^\top \mathbf{d}_k)^2 = \text{Var}[R] \sum_k \|\mathbf{d}_k\|^2 \cos^2 \psi_k = G(1 - S),$$

and subtracting the two mutual informations gives $\frac{1}{2} \log(1 + G) - \frac{1}{2} \log(1 + G(1 - S)) = -\frac{1}{2} \log \frac{1+G-GS}{1+G} = -\frac{1}{2} \log(1 - WS)$. $\square$

Proof of Proposition 3.3. (i) As $\gamma \rightarrow \infty$ we have $w \rightarrow 1$, and (13) gives

$$\ell \rightarrow -\frac{1}{2} \log \cos^2 \psi = -\log \sum_i \sqrt{p_i q_i} = D_B$$

by Proposition 3.2. The convergence is uniform, which is what carries the minimisers with it. Positivity keeps p_τ interior, and $\sqrt{q_i} \geq q_i$ with $\sum_i q_i = 1$ gives $\cos \psi = \sum_i \sqrt{p_i q_i} \geq \min_i \sqrt{p_i}$ for every q in the simplex, so $\cos \psi$ is bounded away from zero uniformly in q and in τ and D_B stays finite. Both objectives are continuous on a compact set and uniformly close on it, so their minimisers converge. (ii) By (14), $\ell = \frac{1}{2} \bar{w} \mathbb{E}_\mu[\psi^2] + O(\psi^4)$ with $d\mu = w d\tau / \bar{w}$; minimising a μ -weighted mean square of a geodesic distance is the definition of the weighted Fréchet mean, and $\psi = \frac{1}{2} d_{\text{FR}}$ converts $\frac{1}{2} \bar{w} \mathbb{E}_\mu[\psi^2]$ into $\frac{1}{8} \bar{w} \mathbb{E}_\mu[d_{\text{FR}}^2]$. $\square$

S3 Two candidate certificates for Proposition 5.1

The C^1 hypothesis cannot be weakened to Lipschitz continuity. A Lipschitz V is differentiable almost everywhere in the state by Rademacher’s theorem, but a single trajectory is a null set and could in principle run inside the exceptional set, so the chain rule step would be unavailable; recovering the argument at that generality needs the Clarke generalised gradient in place of ∇V . Smoothness is also, as it turns out, exactly what the natural candidates lack.

We tried two. The value function of the Dinkelbach problem of Section 5 satisfies (25) with equality in the viscosity sense [1], but value functions of this class are only Lipschitz. Computed on a 41^3 grid and tested on a 61^3 one, (25) fails at about 45% of test points by up to 2.9×10^{-3} . Halving λ moves that violation by 0.6%, so the fault lies in the candidate and not in the bound. At the opposite extreme an affine $V(x) = c_0 - a_0^\top x$ has an exactly constant gradient and makes every requirement linear, so the best certifiable λ is a linear program; over 167,936 constraints it returns $V \equiv 0$ and only the pointwise bound $\min_{x,u} \ell$, which is vacuous here because the reachable floor falls to within sampling resolution of zero at mid-life. Deleting the two boundary requirements lets that program run to any cap imposed on it, which identifies them as what binds: trajectories of (2) increase in every coordinate, so an affine V is monotone along all of them at once and $V(x_0) \geq 0 \geq V|_{\text{failure}}$ forces it to decrease, while the near-aligned states, where the running cost vanishes, demand a local increase.

The two failures are complementary, and together they say what a certificate needs. The value function varies its gradient exactly as the two demands require but is not differentiable where that matters; affine functions are perfectly smooth but cannot vary their gradient at all. A certificate must therefore be both smooth and non-constant in its gradient, which points at a smooth under-approximation of the value function, obtained by sum-of-squares programming when the plant data are polynomial, or by interval arithmetic with an explicit Lipschitz constant.

S4 Numerical verification

S4.1 The identities

The analytical results of the main text were checked numerically before they were used. The rotation identity (20) holds to relative error below 1.1×10^{-7} over the five input–channel pairs $(d, m) = (2, 1), (3, 2), (3, 1), (4, 2)$ and $(5, 3)$; the replicator form to 6.9×10^{-8} ; Theorem 3.6 to 5.1×10^{-6} , limited by finite differencing; and (17) to 1.1×10^{-16} .

S4.2 The exceptional family of Remark 4.5

Integrating the construction of Remark 4.5 for $d = 2$ to 8 with one input holds the misalignment below 1.5×10^{-6} degrees and the mean loss below 10^{-15} . These families carry no prior variance of their own, so their losses are evaluated at a fixed $w = 0.9$ rather than at a γ the plant supplies; the bound has fifteen orders of magnitude of headroom and does not turn on the choice. That is against 1.2×10^{-5} to 5.4×10^{-4} for the best constant input on the same plant.

The family is a knife edge. Perturbing the shape coefficients ν by 1% raises the mean loss to 3.2×10^{-4} under the myopic policy, against 1.8×10^{-4} for the best constant input on the same perturbed plant. At one percent of shape error, steering is already worse than not steering. Closing the loop halves the damage but does not repair it, so the fragility belongs to the construction and not to the open-loop schedule. Mistuning the activation energies instead is tolerated to about two percent.

References

- [1] Martino Bardi and Italo Capuzzo-Dolcetta. *Optimal Control and Viscosity Solutions of Hamilton–Jacobi–Bellman Equations*. Birkhäuser, 1997. doi: 10.1007/978-0-8176-4755-1.